


+ . \* 2021 SMARCLE Paper Review . \* +

# Transformer - Attention Is All You Need (NIPS 2017)

지능기전공학부 무인이동체공학전공  
송혜원

# Seq2Seq의 문제점



# Attention


이전엔 입력을  
하나의 **context vector**로 압축



**Attention** 기법은  
문장 자체를 입력으로!



어떤 단어가 어떤 단어와  
가장 연관성이 있는지 알 수 있게 됨



# Transformer<sub>Input</sub>




Figure 1: The Transformer - model architecture.

Dog loves potato

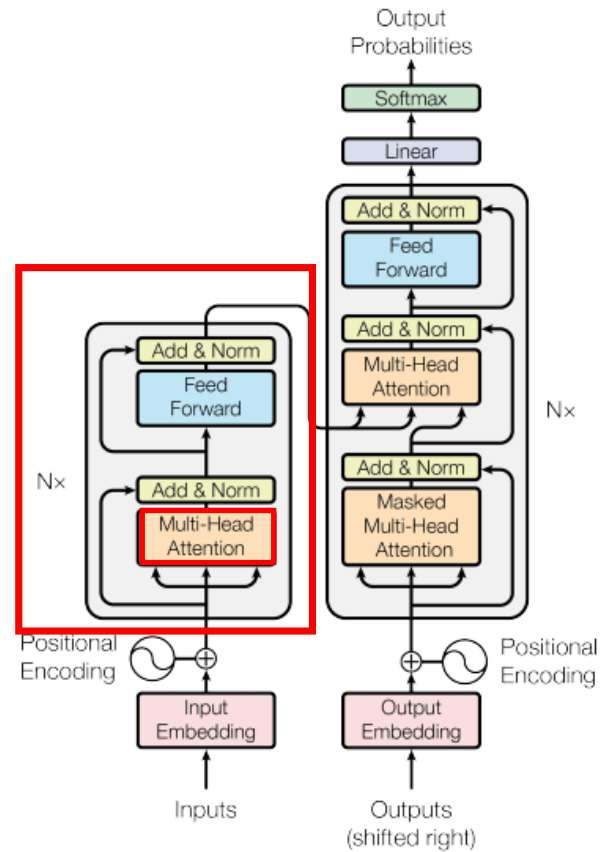
Dog			...				
loves			...				
potato			...				

512

문장 전체를 통째로 넣다보니 단어의 위치 정보를 알 수 없음

➡ **Positional Encoding** 통해 위치 정보를 따로 더해줌!

# Transformer\_Encoder Self-Attention



Encoder Self-Attention:




Figure 1: The Transformer - model architecture.

# Transformer\_Encoder

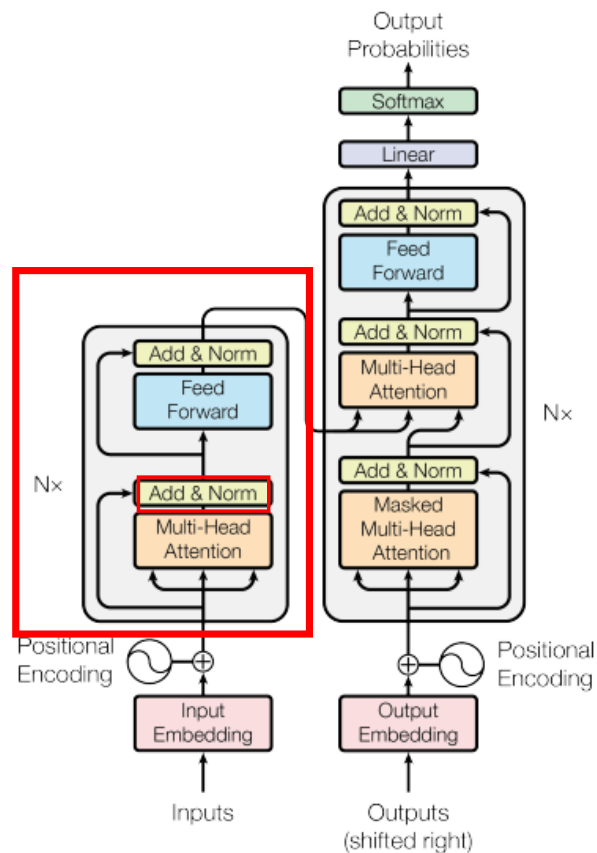



Figure 1: The Transformer - model architecture.

Encoder Self-Attention:



Residual learning

➡ 특정 layer를 건너뛸

Gradient vanishing 완화

Global optima 잘 찾을 수 있음

# Transformer\_Encoder

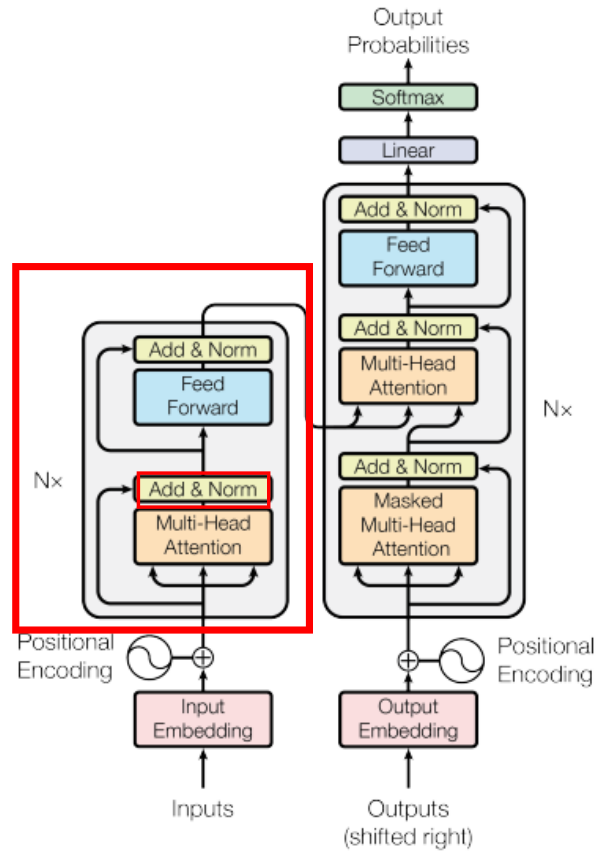



Figure 1: The Transformer - model architecture.

Encoder Self-Attention:



Residual learning

➡ 특정 layer를 건너뛸

Gradient vanishing 완화

Global optima 잘 찾을 수 있음

인풋의 다양한 특징 **attention** 위해 n개의 **layer** 병렬로 쌓음

# Transformer\_Masked Decoder Self-Attention

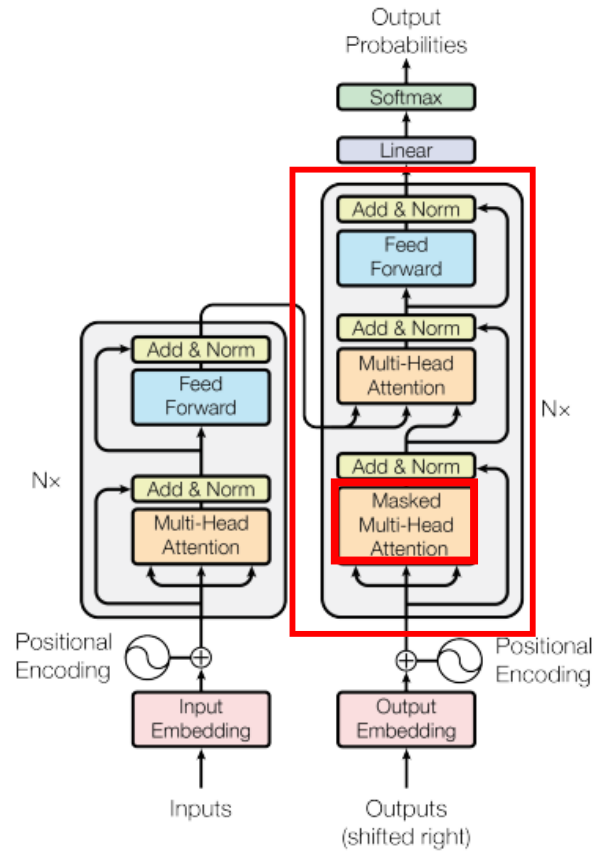


Figure 1: The Transformer - model architecture.

Masked Decoder Self-Attention:

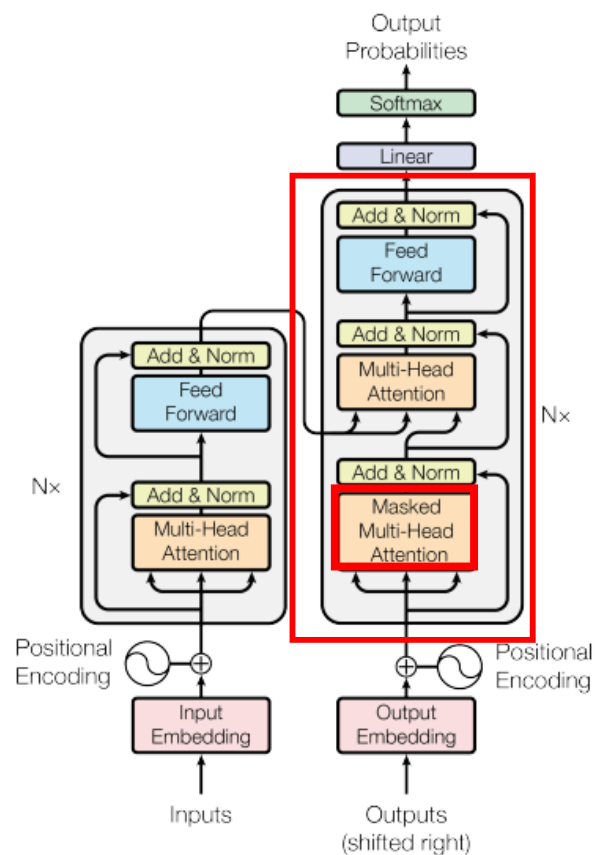
강아지는 감자를 좋아해



The diagram shows the words '강아지는', '감자를', and '좋아해' in green boxes. Arrows indicate the attention mechanism: a blue arrow from '강아지는' to '감자를', an orange arrow from '감자를' to '좋아해', and a red arrow from '강아지는' to '좋아해'.



# Transformer\_Masked Decoder Self-Attention



Masked Decoder Self-Attention:



Encoder Self-Attention:




Figure 1: The Transformer - model architecture.

# Transformer\_Masked Decoder Self-Attention

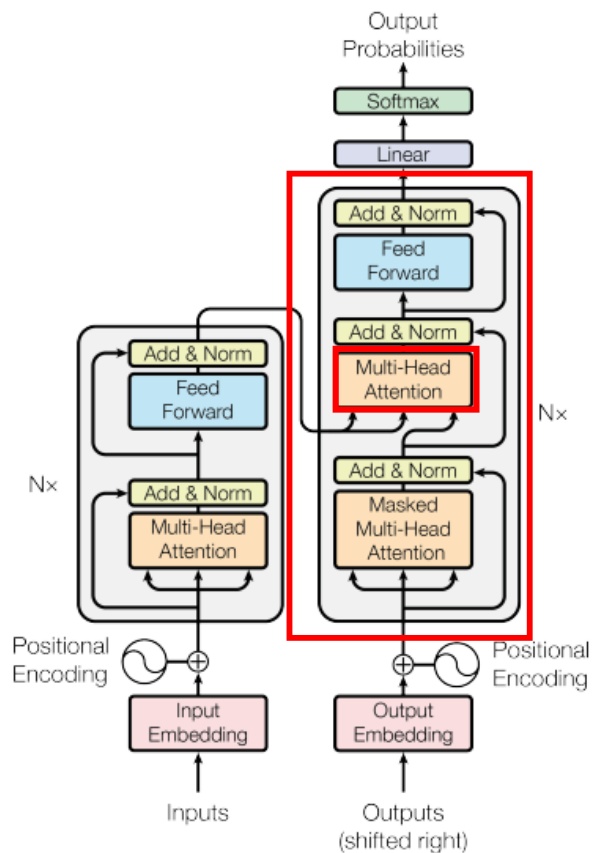


Figure 1: The Transformer - model architecture.

Encoder-Decoder Attention:

Dog loves potato 강아지

- 인코더 파트에서 나온 출력 결과를 디코더마다 적용하여 전적으로 활용
- 시퀀스가 끝날 때까지 반복

# Transformer\_Multi-Head Attention?

Q(Query) : 물어보는 주체    K(Key) : 대상    V(Value) : 가중치

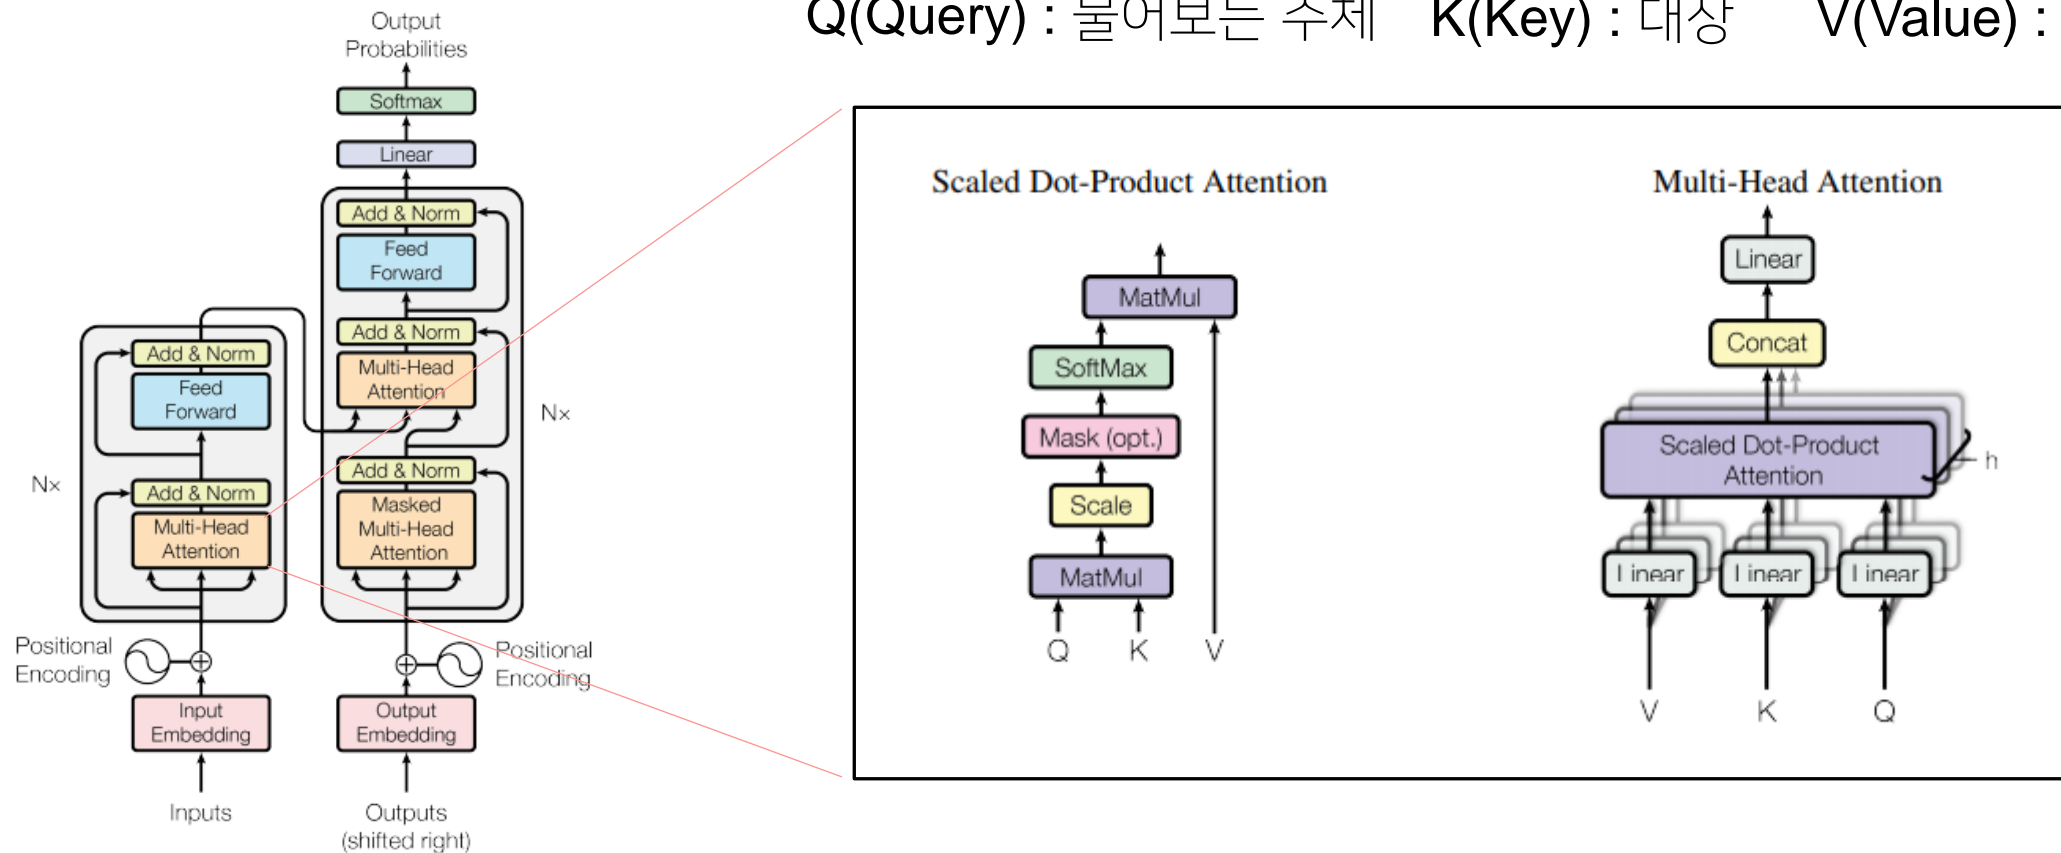


Figure 1: The Transformer - model architecture.

# Transformer\_Multi-Head Attention?

Q(Query) : 물어보는 주체    K(Key) : 대상    V(Value) : 가중치

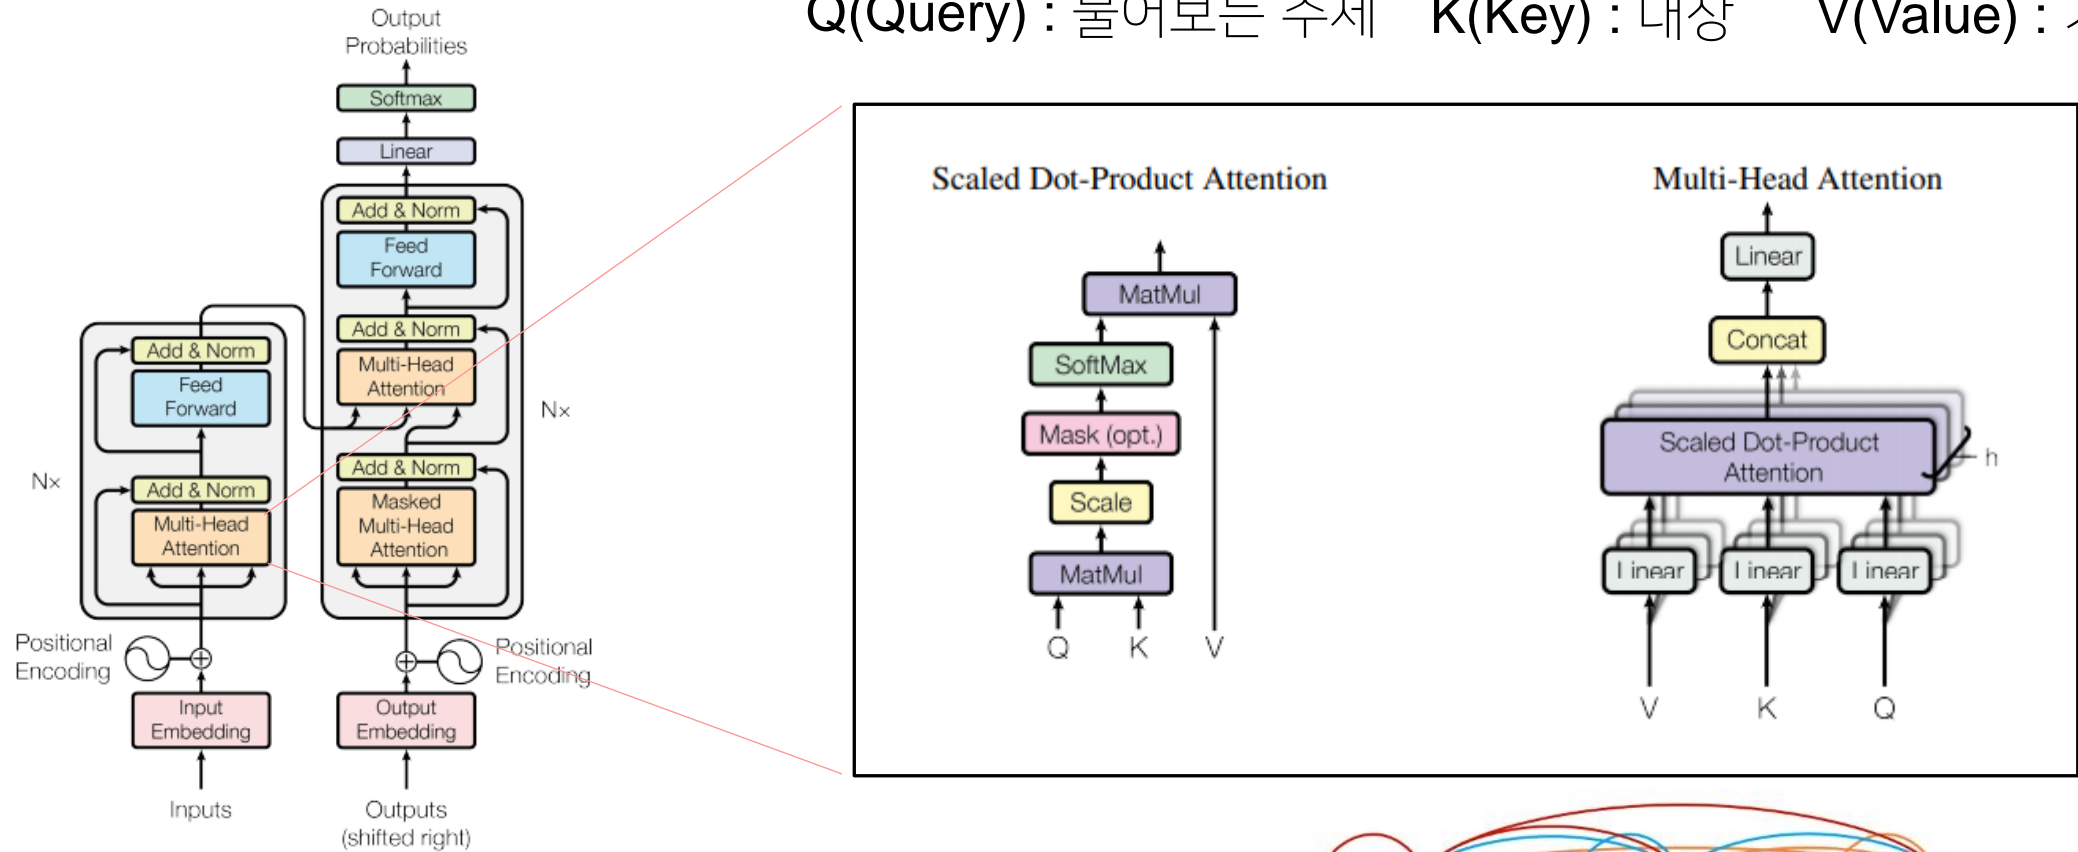



Figure 1: The Transformer - model architecture.

Encoder Self-Attention:



# Transformer\_Multi-Head Attention?

Q(Query) : 물어보는 주체    K(Key) : 대상    V(Value) : 가중치


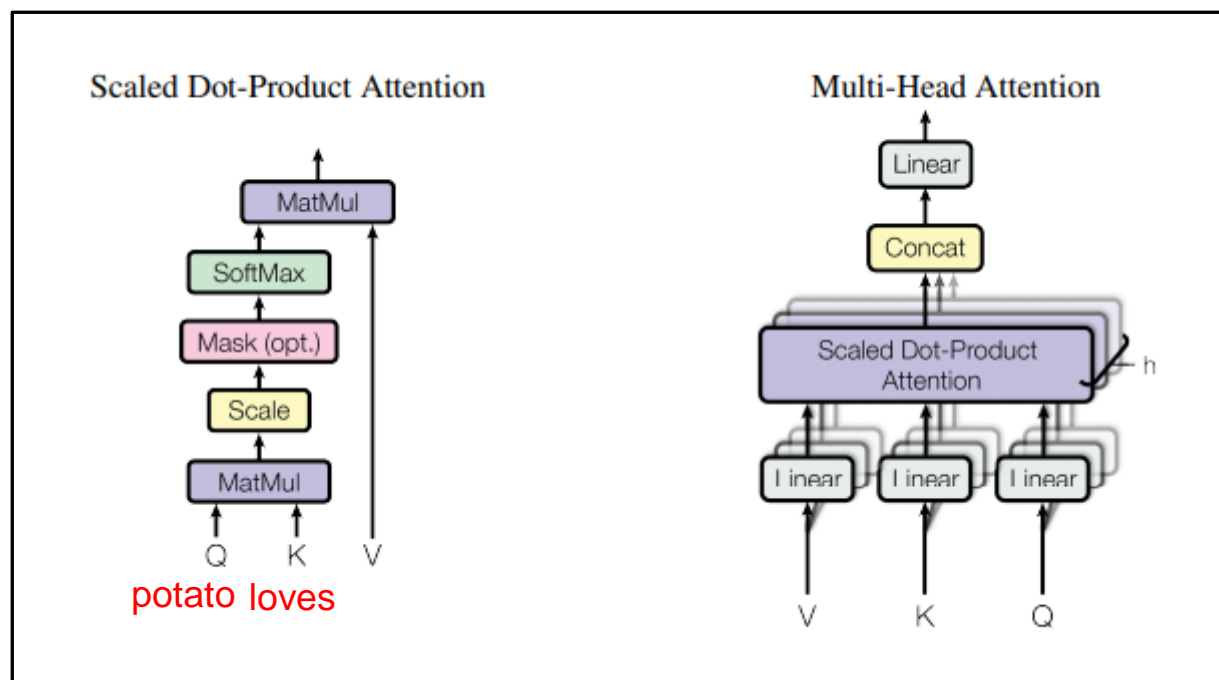



Figure 1: The Transformer - model architecture.

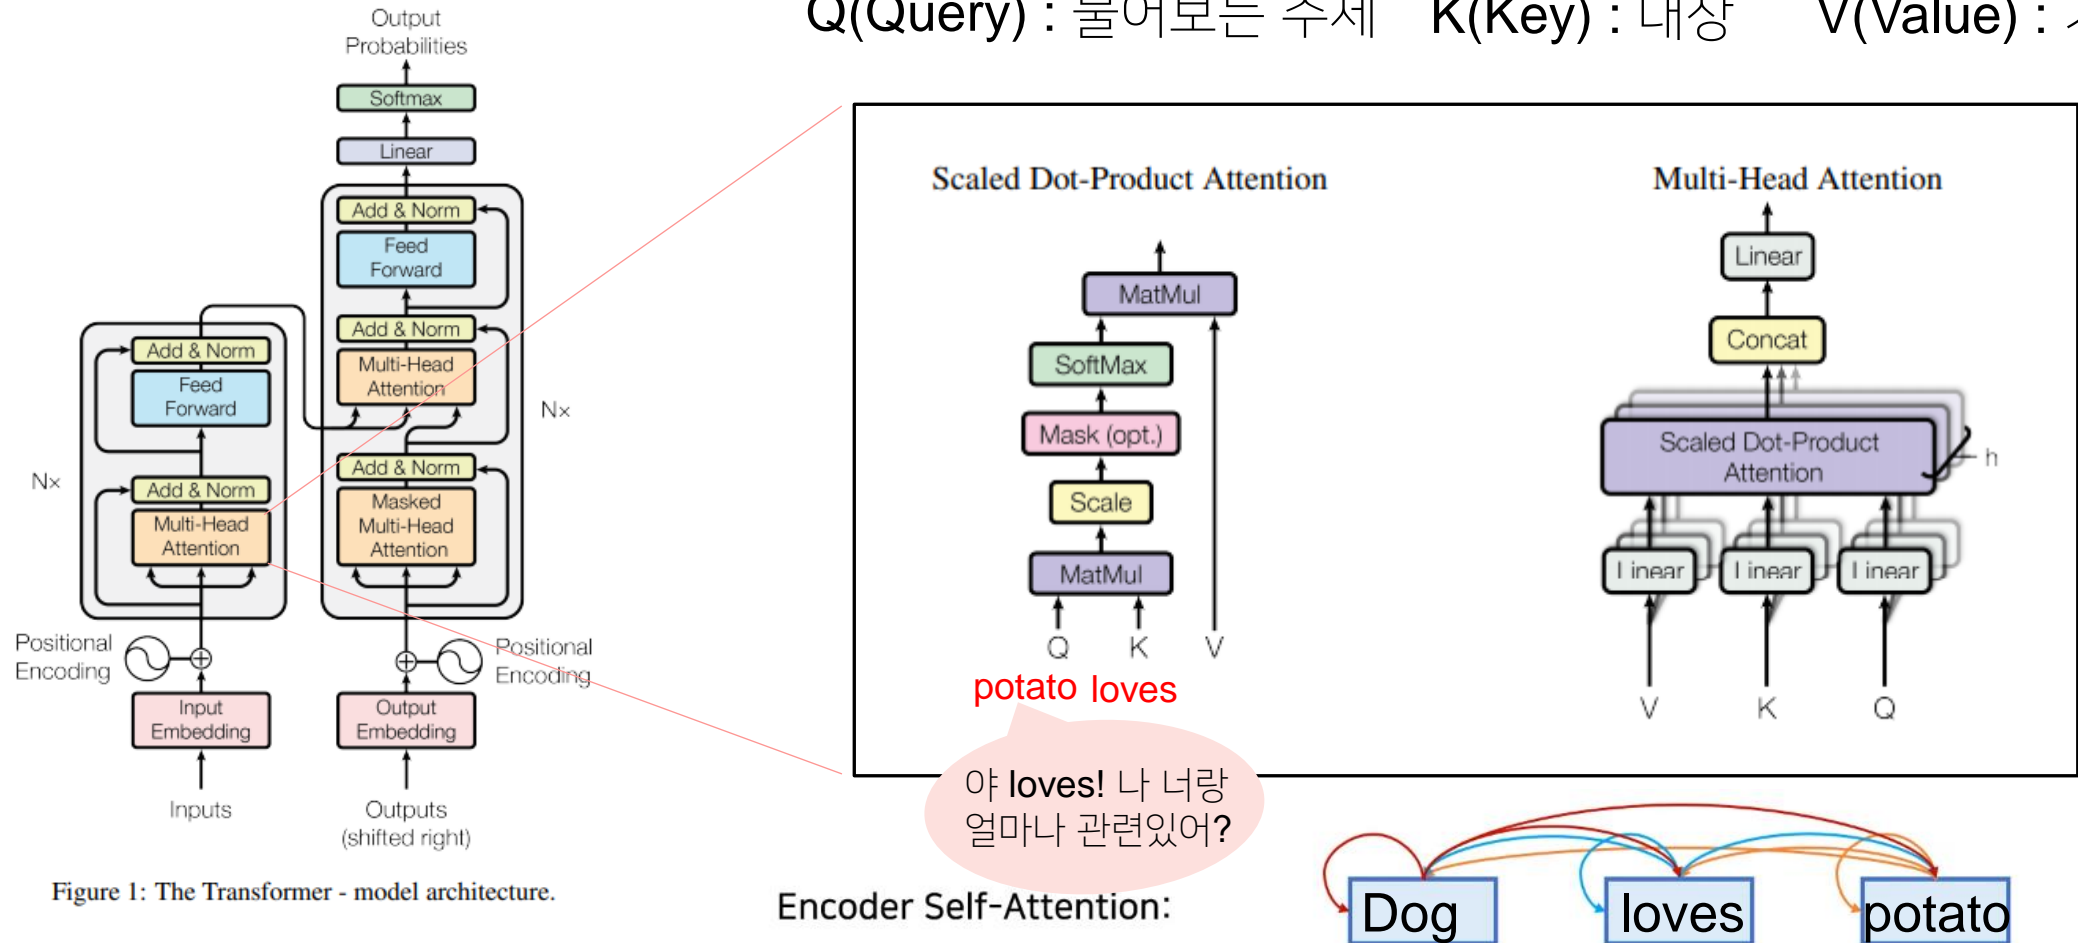


Encoder Self-Attention:



# Transformer\_Multi-Head Attention?

Q(Query) : 물어보는 주체    K(Key) : 대상    V(Value) : 가중치



# Transformer\_Multi-Head Attention?

Q(Query) : 물어보는 주체    K(Key) : 대상    V(Value) : 가중치

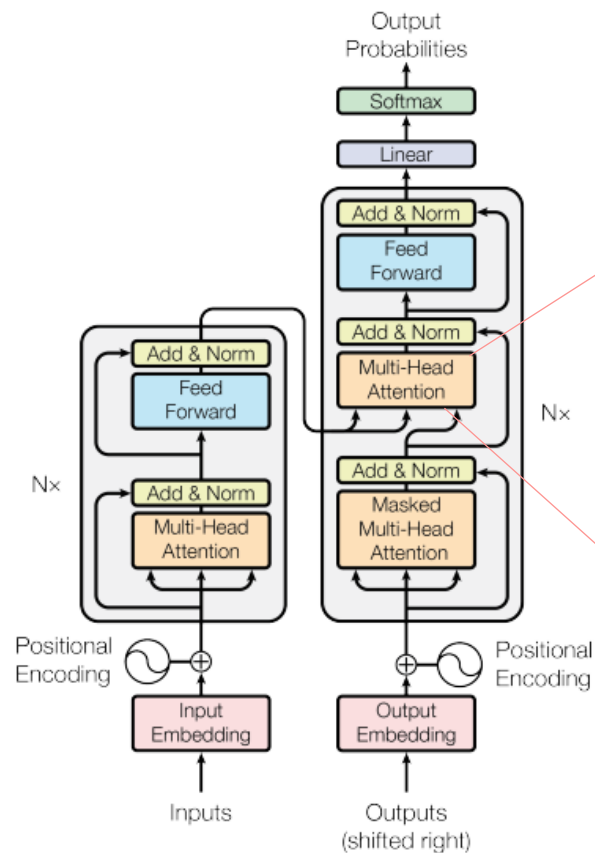
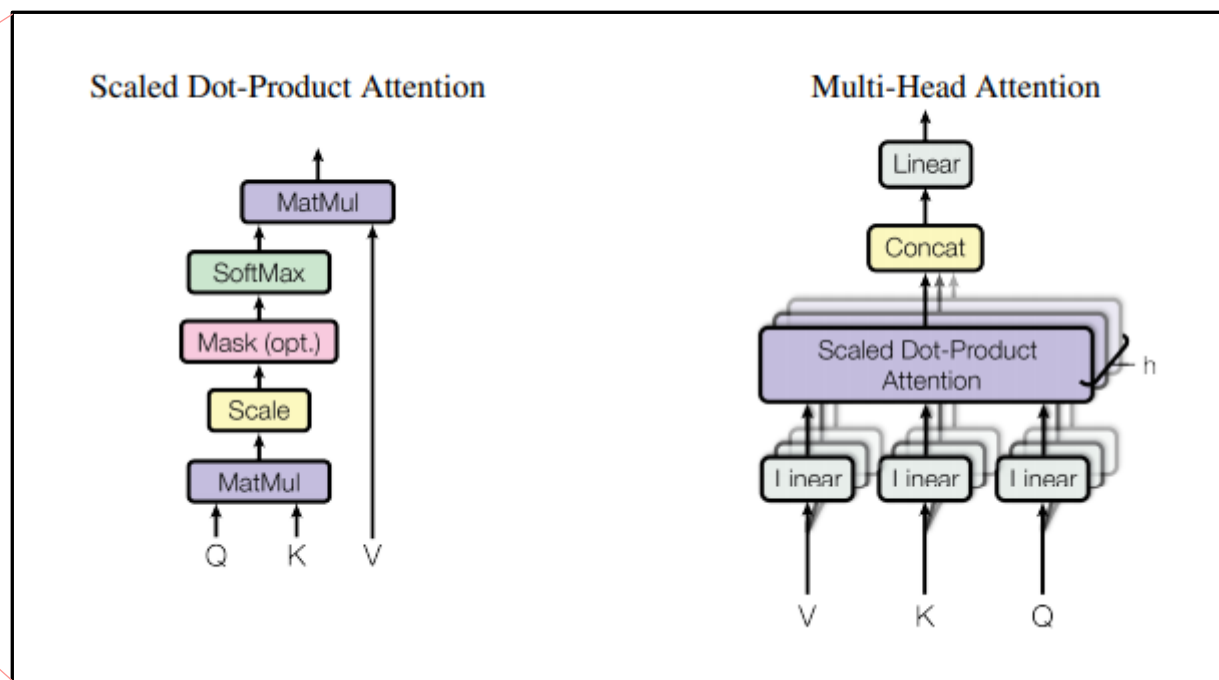


Figure 1: The Transformer - model architecture.



Encoder-Decoder Attention:

Dog

loves

potato

강아지

# Transformer\_Multi-Head Attention?

Q(Query) : 물어보는 주체    K(Key) : 대상    V(Value) : 가중치

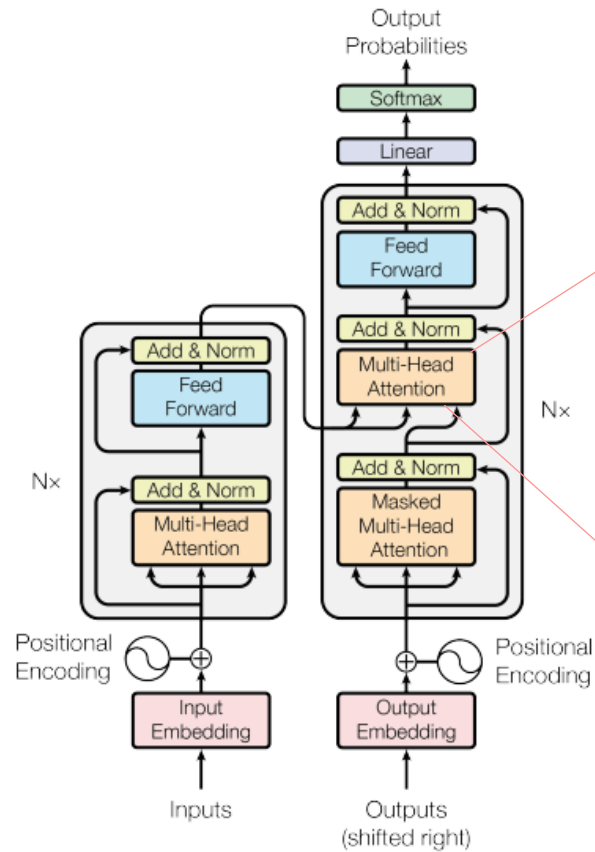
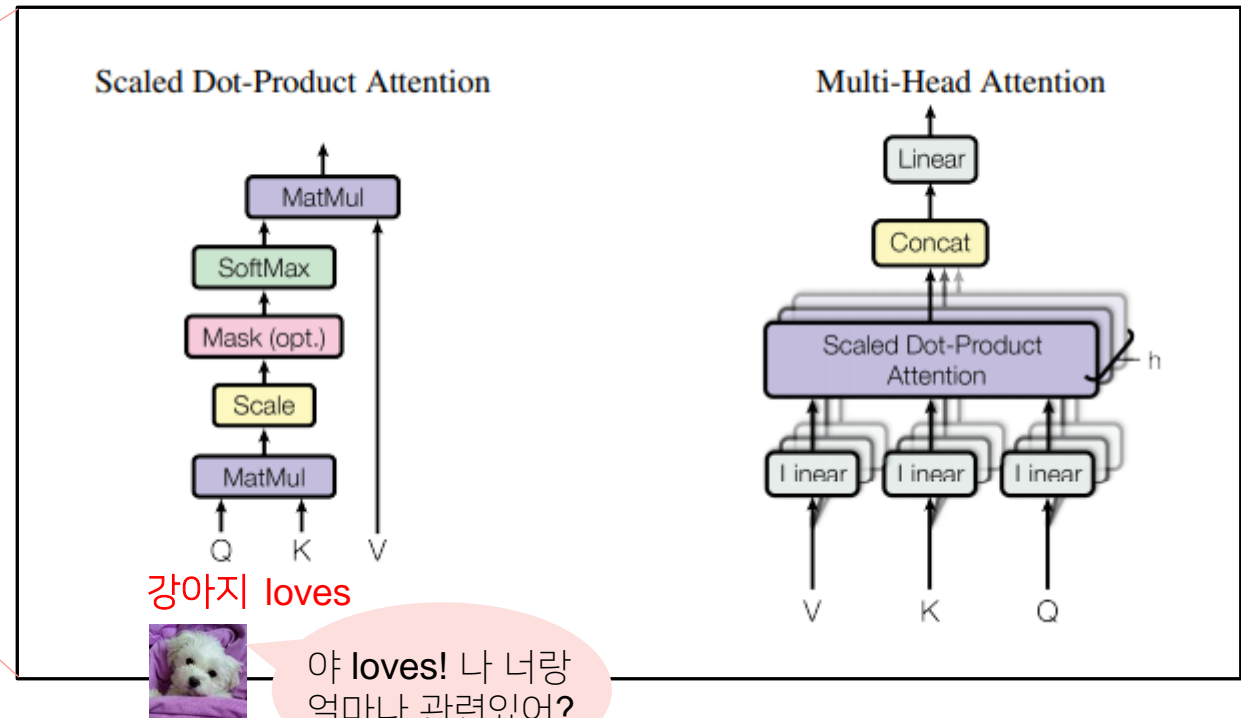


Figure 1: The Transformer - model architecture.



Encoder-Decoder Attention:

Dog

loves

potato

강아지



# Transformer\_Multi-Head Attention?

Q(Query) : 물어보는 주체    K(Key) : 대상    V(Value) : 가중치


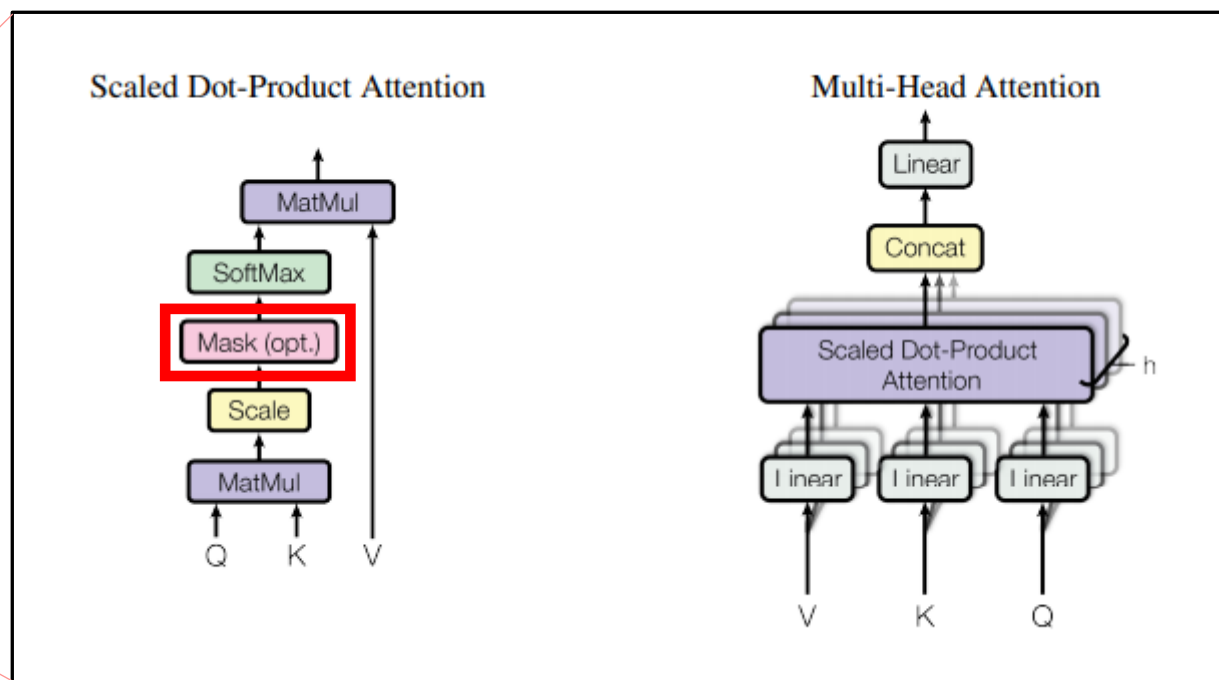


Figure 1: The Transformer - model architecture.



Masked Decoder Self-Attention:

강아지는 감자를 좋아해

# Transformer\_Multi-Head Attention?

Q(Query) : 물어보는 주체    K(Key) : 대상    V(Value) : 가중치


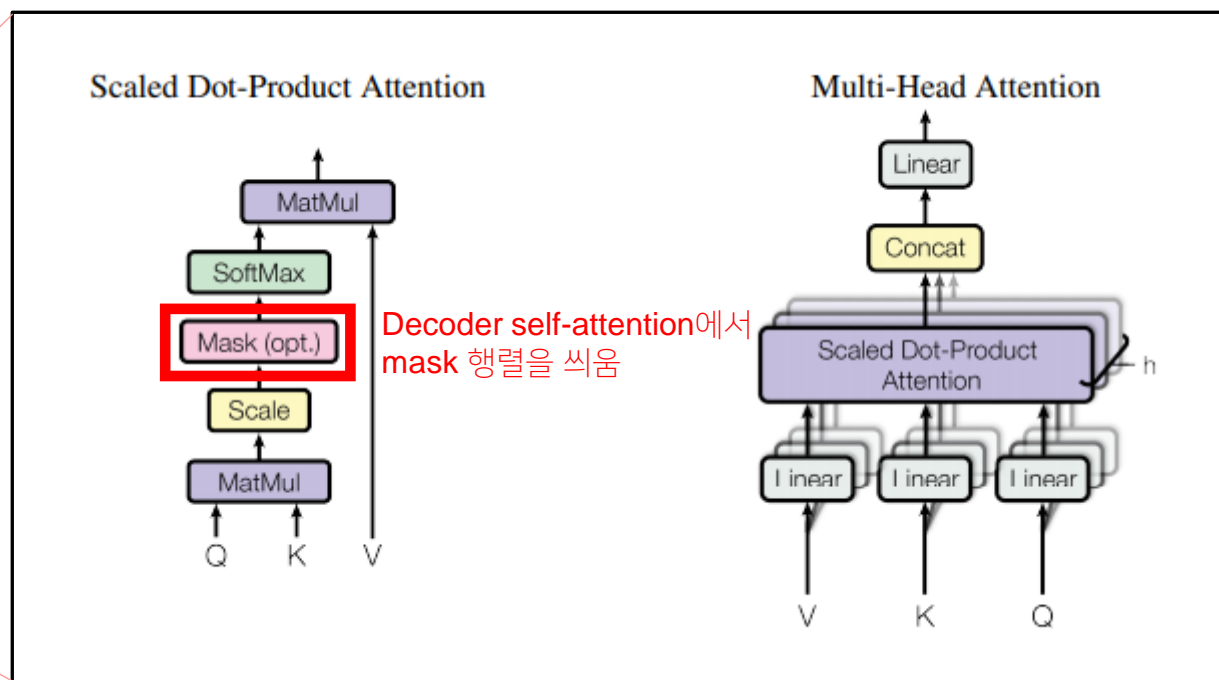


Figure 1: The Transformer - model architecture.



Masked Decoder Self-Attention:

강아지는 감자를 좋아해

# Transformer\_Positional Encoding




Figure 1: The Transformer - model architecture.

In this work, we use sine and cosine functions of different frequencies:

$$PE_{(pos, 2i)} = \sin(pos/10000^{2i/d_{model}})$$

$$PE_{(pos, 2i+1)} = \cos(pos/10000^{2i/d_{model}})$$

where  $pos$  is the position and  $i$  is the dimension. That is, each dimension of the positional encoding corresponds to a sinusoid. The wavelengths form a geometric progression from  $2\pi$  to  $10000 \cdot 2\pi$ . We chose this function because we hypothesized it would allow the model to easily learn to attend by relative positions, since for any fixed offset  $k$ ,  $PE_{pos+k}$  can be represented as a linear function of  $PE_{pos}$ .

We also experimented with using learned positional embeddings [9] instead, and found that the two versions produced nearly identical results (see Table 3 row (E)). We chose the sinusoidal version because it may allow the model to extrapolate to sequence lengths longer than the ones encountered during training.

- 본문에서는 주기함수를 이용해서 단어의 위치를 **정해줌**
- 이후 모델들은 **학습이 가능한** 임베딩 레이어를 넣어주는 추세

# Transformer\_Conclusion

Table 3: Variations on the Transformer architecture. Unlisted values are identical to those of the base model. All metrics are on the English-to-German translation development set, newstest2013. Listed perplexities are per-wordpiece, according to our byte-pair encoding, and should not be compared to per-word perplexities.

		$N$	$d_{\text{model}}$	$d_{\text{ff}}$	$h$	$d_k$	$d_v$	$P_{\text{drop}}$	$\epsilon_{ls}$	train steps	PPL (dev)	BLEU (dev)	params $\times 10^6$
기본	base	6	512	2048	8	64	64	0.1	0.1	100K	4.92	25.8	65
	(A)				1	512	512				5.29	24.9	
					4	128	128				5.00	25.5	
					16	32	32				4.91	25.8	
				32	16	16				5.01	25.4		
key dimension	(B)				16						5.16	25.1	58
					32						5.01	25.4	60
model 크기	(C)	2									6.11	23.7	36
		4									5.19	25.3	50
		8									4.88	25.5	80
			256			32	32				5.75	24.5	28
			1024			128	128				4.66	26.0	168
				1024							5.12	25.4	53
				4096							4.75	26.2	90
dropout	(D)							0.0			5.77	24.6	
								0.2			4.95	25.5	
									0.0		4.67	25.3	
									0.2		5.47	25.7	
p.e. 학습	(E)	positional embedding instead of sinusoids									4.92	25.7	
	big	6	1024	4096	16			0.3		300K	<b>4.33</b>	<b>26.4</b>	213

# Transformer\_TBC..



감사합니다 :-3